

SIMATS ENGINEERING

ASSESSMENT TOOL 1

COURSE NAME: Artificial Intelligence

COURSE CODE: CSA17

COURSE OUTCOME COVERED

C05: *Design AI-based systems using PySpark, RDD operations, and intelligent agent architectures, using scalable systems and integrate components in distributed AI applications. (BTL 5)*

Assessment Tool Weightage: 50%

QUESTIONS: 5

Annexure A

Duration : 1 Hour
Total Marks:50

Design Task

Code of Conduct:

I Arava Raja (192365033) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.

I understand that any violation of academic integrity rules will result in disciplinary action.

Arava Raja

Signature of the student

Design Task

1. Hybrid AI Recommendation Engine Design

Design a hybrid AI-based recommendation system that integrates content-based and collaborative filtering techniques using PySpark RDDs. Explain how distributed data processing can be used to handle large-scale user-item interactions, and describe the role of intelligent agents in adapting recommendations based on user behavior and feedback in real time.

2. Smart Disaster Detection System

Develop a distributed AI system using PySpark and RDD operations to detect natural disasters (such as floods or earthquakes) from multi-source data streams (satellite images, sensors, and social media). Describe the system architecture, data fusion techniques, and how intelligent agents coordinate to provide early warnings and response strategies.

3. AI-Based Fraud Detection Framework

Design a scalable fraud detection system using PySpark RDDs to process large volumes of financial transactions. Explain how anomaly detection algorithms can be implemented in a distributed manner and how intelligent agents collaborate to identify suspicious patterns, trigger alerts, and continuously improve detection accuracy.

4. Autonomous Supply Chain Optimization System

Construct a distributed AI system using PySpark for optimizing supply chain operations (inventory, logistics, demand forecasting). Describe how RDD transformations support large-scale data processing and how multiple intelligent agents interact to make decentralized decisions for cost minimization and efficiency improvement.

5. Personalized Learning Analytics Platform

Design an AI-driven educational platform that uses PySpark and RDDs to analyze student learning behavior and performance data. Explain how the system supports adaptive learning through intelligent agents, real-time feedback, and scalable analytics to personalize content delivery for thousands of learners.

RUBRICS FOR CALCULATION

Evaluation Criteria	Problem Understanding & Requirement Analysis	System Architecture Design	Use of PySpark & RDD Operations	Intelligent Agent Integration	Scalability & Distributed Computing Design	Data Flow & Processing Pipeline	Innovation & Creativity	Clarity, Presentation & Documentation
Weightage	10%	20%	15%	15%	10%	10%	10%	10%

Criteria	Excellent (5)	Good (4)	Average (3)	Poor (2)
Problem Understanding & Requirement Analysis	Clearly defines problem, objectives, and constraints	Mostly clear with minor gaps	Basic understanding	Unclear or incorrect
System Architecture Design	Complete, scalable, well-structured architecture with diagrams	Good design with minor issues	Basic design, lacks detail	Poor/incomplete
Use of PySpark & RDD Operations	Efficient and correct use of RDD transformations/actions	Mostly correct usage	Limited or partial use	Incorrect/missing
Intelligent Agent Integration	Strong integration with clear agent roles and logic	Good integration	Basic integration	Poor/no integration
Scalability & Distributed Computing Design	Clearly addresses scalability, fault tolerance, parallelism	Covers most aspects	Limited consideration	Not addressed
Data Flow & Processing Pipeline	Clear, logical, and well-defined data flow	Mostly clear	Some gaps	Poorly defined

Innovation & Creativity	Highly innovative, realistic, and impactful solution	Some innovation	Basic idea	No originality
Clarity, Presentation & Documentation	Very clear, neat diagrams, well-organized report	Mostly clear	Somewhat organized	Poor presentation

ARAVARAJA(192365033)

1. Hybrid AI Recommendation Engine Design

Introduction


A Hybrid AI Recommendation Engine is an intelligent recommendation system that combines **content-based filtering** and **collaborative filtering** to provide personalized recommendations. The system uses **PySpark RDDs** to process large-scale user-item interaction data in a distributed environment.


The system analyzes user preferences, previous purchases, ratings, browsing history, and item characteristics. Intelligent agents continuously monitor user behavior and feedback to update recommendations in real time.

Objectives

- Provide personalized recommendations.
- Process millions of user-item interactions efficiently.
- Combine content-based and collaborative filtering.
- Adapt recommendations based on real-time feedback.
- Improve recommendation accuracy.
- Support thousands or millions of users.

System Architecture





Content-Based Filtering

Content-based filtering recommends items similar to items previously liked by the user.

For example, if a user frequently watches action movies, the system recommends other action movies.

Features may include:

- Category
- Genre
- Price
- Brand
- Keywords
- Item description
- User interests

A feature vector is created for every item, and similarity between items can be calculated using **Cosine Similarity**.

Collaborative Filtering

Collaborative filtering uses the behavior of similar users.

For example:

User A and User B have purchased similar products. User A purchased Product X, but User B has not. Product X can therefore be recommended to User B.

The system can construct:

User → Item → Rating

and identify similar users or similar items.

PySpark RDD Operations

The interaction data can be represented as:

(userID, itemID, rating)

RDD transformations can include:

`ratings.map(lambda x: (x[0], x[2]))`

to extract user-rating information.

Other useful operations include:

- map()
- filter()
- flatMap()
- groupByKey()
- reduceByKey()
- join()
- mapValues()

Actions include:

- count()
- collect()
- take()
- reduce()

For example, reduceByKey() can calculate aggregate ratings efficiently across distributed partitions.

Intelligent Agents

Three major agents can be used:

1. User Behavior Agent

Monitors:

- Searches
- Clicks
- Purchases
- Ratings
- Watch time

2. Recommendation Agent

Combines content-based and collaborative results and generates recommendations.

3. Feedback Agent

Analyzes whether the user accepted, ignored, clicked, or purchased recommended items.

The feedback is then sent back to the recommendation engine.

Real-Time Adaptation

The system continuously updates user profiles.

For example:

User searches for laptops



Agent detects new interest



User preference updated



Recommendation model updated



New laptop recommendations

Scalability

PySpark distributes RDD partitions across multiple worker nodes. Therefore, millions of interactions can be processed in parallel.

Advantages include:

- Parallel processing
- Fault tolerance
- Horizontal scalability
- Faster computation
- Distributed memory processing

Conclusion

The proposed hybrid recommendation engine combines the advantages of content-based and collaborative filtering. PySpark provides scalable data processing, while intelligent agents continuously monitor user behavior and feedback. This produces accurate, adaptive, and personalized recommendations.

2. Smart Disaster Detection System

Introduction

A Smart Disaster Detection System is a distributed AI system designed to detect natural disasters such as **floods and earthquakes** using multiple data sources including satellite images, sensors, and social media. The assessment specifically requires PySpark/RDD processing, data fusion, and intelligent-agent coordination.

Objectives

- Detect disasters at an early stage.
- Process data from multiple sources.
- Identify abnormal environmental conditions.
- Generate early warnings.
- Reduce disaster response time.
- Coordinate emergency response.

System Architecture

Satellite Images



Sensors ----->|

v

Social Media ----->|

+-----+

| Data |

| Collection |

+-----+

|

v

+-----+

| PySpark RDD |

| Processing |

+-----+

|


v

+-----+

| Data Fusion |

+-----+

|



Data Sources

Satellite Data

Satellite images can provide information about:

- Flooded regions
- Cloud patterns
- Water-level changes
- Land movement

Sensor Data

Sensors can provide:

- Water levels
- Temperature
- Pressure
- Ground vibrations
- Rainfall

Social Media

Social media can provide real-time information such as:

- User reports
- Emergency messages
- Location information
- Disaster-related keywords

PySpark RDD Processing

Each data source can be converted into RDDs.

For example:

Sensor RDD

(location, water_level, timestamp)

Social Media RDD

(location, message, timestamp)

RDD transformations such as:

- map()
- filter()
- flatMap()
- reduceByKey()
- join()

can be used to process the data.

For example, sensor readings above a predefined threshold can be filtered:

`sensorRDD.filter(lambda x: x[1] > threshold)`

Data Fusion

Data fusion combines information from different sources.

For example:

High rainfall

+

High water level

+

Satellite detects flooded area

+

Social media reports flooding

↓

High-confidence flood detection

Using multiple sources reduces false alarms.

Intelligent Agents

Sensor Monitoring Agent

Continuously monitors sensor values.

Satellite Analysis Agent

Analyzes satellite information and detects environmental changes.

Social Media Agent

Identifies disaster-related messages and extracts locations.

Decision Agent

Combines information from all agents and determines disaster severity.

Emergency Response Agent

Sends alerts and recommends suitable response actions.

Early Warning Process

Data Collection

↓

RDD Processing

↓

Anomaly Detection

↓

Data Fusion

↓

Disaster Confirmation

↓

Severity Estimation

↓
Warning Generation

↓
Emergency Response

Scalability

PySpark allows data to be divided into multiple partitions and processed simultaneously. This is useful because disaster systems may receive thousands of sensor readings and social-media messages simultaneously.

Conclusion

The Smart Disaster Detection System provides a scalable and intelligent solution for early disaster detection. Combining multiple data sources with PySpark distributed processing and intelligent agents improves detection accuracy, reduces response time, and supports real-time emergency management.

3. AI-Based Fraud Detection Framework

Introduction

An AI-Based Fraud Detection Framework identifies suspicious financial transactions by processing large volumes of transaction data using PySpark RDDs. The assessment requires distributed anomaly detection and intelligent agents that collaborate to identify suspicious patterns and continuously improve detection accuracy.

Objectives

- Detect fraudulent transactions.
- Process large transaction datasets.
- Identify abnormal behavior.
- Generate real-time alerts.
- Reduce false positives.
- Continuously improve fraud detection.

Architecture

Financial Transactions

|
v

Data Collection

|
v

PySpark RDD

|
v

Data Cleaning


|
v

Feature Extraction

|
v

Anomaly Detection

|



Transaction Data

Each transaction may contain:

Transaction ID

User ID

Amount

Location

Timestamp

Merchant

Payment Method

Device ID

RDD Processing

Transaction records can be represented as:

(transactionID, userID, amount, location, timestamp)

RDD operations can clean and transform the data.

Example:

`transactions.filter(lambda x: x.amount > 0)`

Transactions can be grouped by user:

`transactions.map(lambda x: (x.userID, x)).groupByKey()`

`reduceByKey()` can be used to calculate transaction totals.

Feature Extraction

Important fraud detection features include:

- Transaction amount
- Transaction frequency
- Geographic distance
- Time of transaction
- Device changes
- Unusual merchant
- Previous transaction history

Anomaly Detection

A transaction is considered suspicious if it significantly differs from the user's normal behavior.

Example:

Normal:

User usually spends ₹500–₹2,000

New transaction:

₹1,50,000 from another country



High anomaly score



Potential fraud

Machine-learning algorithms such as clustering or anomaly-detection models can be applied to distributed data.

Intelligent Agents

Transaction Monitoring Agent

Monitors incoming transactions.

User Behavior Agent

Maintains the user's normal spending pattern.

Anomaly Detection Agent

Calculates anomaly scores.

Risk Assessment Agent

Assigns a fraud-risk score.

Alert Agent

Generates alerts when risk exceeds a threshold.

Agent Collaboration

Transaction Agent



Behavior Agent



Anomaly Agent



Risk Agent



Alert Agent

Continuous Learning

When investigators classify a transaction as fraud or legitimate, the result is used as feedback.

Transaction



Prediction



Fraud/Normal



Human Feedback



This helps improve future detection accuracy.

PySpark distributes transactions across cluster nodes. Therefore, millions of transactions can be processed in parallel.

- Fault tolerance
- Parallel computation
- Distributed storage
- High throughput
- Horizontal scalability


The proposed framework combines PySpark RDD processing, anomaly detection, and intelligent agents to detect fraudulent financial transactions efficiently. Continuous feedback allows the system to improve its accuracy over time.

Introduction


An Autonomous Supply Chain Optimization System uses distributed AI to optimize **inventory, logistics, and demand forecasting**. PySpark RDD transformations process large-scale supply-chain data, while multiple intelligent agents make decentralized decisions for cost minimization and efficiency improvement.

- Optimize inventory levels.
- Predict product demand.
- Select efficient transportation routes.
- Minimize logistics costs.
- Avoid stock-outs and overstocking.
- Make decentralized supply-chain decisions.


Sales Data Supplier Data



Demand Forecast Inventory Analysis




Intelligent Agents



Inventory Logistics Supplier

Agent Agent Agent



Optimized Decision

Data Sources

The system processes:

- Historical sales
- Inventory records
- Supplier information
- Transportation costs
- Delivery times
- Customer orders
- Seasonal demand

RDD Transformations

Data can be represented as:

(productID, quantity, date, location)

RDD transformations include:

map()

Transforms raw records into useful structures.

filter()

Removes invalid or unnecessary records.

reduceByKey()

Calculates total product sales.

groupByKey()

Groups data by product or location.

join()

Combines inventory and sales information.

Example:

```
sales.map(lambda x: (x.productID, x.quantity))  
      .reduceByKey(lambda a, b: a + b)
```

Demand Forecasting

Historical sales are analyzed to estimate future demand.

Historical Sales



RDD Processing



Pattern Analysis



Demand Forecast



Inventory Decision

If demand is expected to increase, the system can recommend increasing inventory.

Intelligent Agents

Inventory Agent

Monitors stock levels and recommends replenishment.

Logistics Agent

Selects efficient delivery routes.

Supplier Agent

Evaluates suppliers based on:

- Cost
- Quality
- Delivery time
- Reliability

Demand Forecasting Agent

Predicts future product demand.

Optimization Agent

Combines the outputs of other agents and selects the best decision.

Decentralized Decision Making

Agents can operate independently but share information.

For example:

Demand Agent → predicts high demand



Inventory Agent → requests additional stock



Supplier Agent → selects supplier



Logistics Agent → selects delivery route



Optimization Agent → final decision

Cost Minimization

The system can consider:

Total Cost =

Inventory Cost

+ Transportation Cost

+ Ordering Cost

+ Storage Cost

The optimal solution minimizes the overall cost while maintaining sufficient inventory.

Scalability

PySpark allows large supply-chain datasets to be processed across multiple machines.

This enables the system to handle:

- Multiple warehouses
- Thousands of products
- Large numbers of suppliers
- Millions of transactions

Conclusion

The autonomous supply-chain system combines distributed PySpark processing with intelligent agents. It enables organizations to forecast demand, optimize inventory, select efficient routes, and reduce operational costs.

5. Personalized Learning Analytics Platform

Introduction


A Personalized Learning Analytics Platform is an AI-driven educational system that analyzes student learning behavior and performance using PySpark and RDDs. Intelligent agents provide adaptive learning, real-time feedback, and personalized content delivery for thousands of learners.


Objectives

- Analyze student performance.
- Identify strengths and weaknesses.
- Recommend personalized learning materials.
- Provide real-time feedback.
- Detect students requiring additional support.
- Scale the platform for thousands of students.

System Architecture

Student Data





Student Data

The system can collect:

- Test scores
- Assignment marks
- Attendance
- Time spent studying
- Video watching time
- Quiz attempts
- Number of incorrect answers
- Learning progress

RDD Processing

Student records can be represented as:

(studentID, subject, score, studyTime)

RDD operations can process this data.

Example:

`studentRDD.filter(lambda x: x.score < 50)`

can identify students with low scores.

`map()` can transform records, while `reduceByKey()` can calculate average performance.

Learning Analytics

The system calculates:

- Average score
- Subject-wise performance

- Learning speed
- Weak topics
- Strong topics
- Engagement level

Example:

Mathematics → 85%

Programming → 72%

Cybersecurity → 45%

Weak Area → Cybersecurity

The system can recommend additional cybersecurity learning materials.

Intelligent Agents

Student Monitoring Agent

Tracks student activity and performance.

Performance Analysis Agent

Identifies strengths and weaknesses.

Content Recommendation Agent

Selects suitable learning materials.

Feedback Agent

Provides immediate feedback after tests and activities.

Intervention Agent

Identifies students who require additional assistance.

Adaptive Learning

The platform adapts content according to the student's performance.

Student completes quiz



Performance analyzed



Weak topic identified



Difficulty adjusted



Personalized content provided



Student completes activity



Performance updated

For a high-performing student, the system can provide advanced questions. For a struggling student, it can provide simpler explanations and additional practice.

Real-Time Feedback

Suppose a student answers five questions incorrectly on a particular topic.

The system can immediately:

1. Identify the weak concept.
2. Provide an explanation.

3. Recommend a tutorial.
4. Generate additional practice questions.
5. Monitor improvement.

Scalability

PySpark enables the platform to process learning data from thousands of students simultaneously.

RDD partitioning provides:

- Parallel processing
- Distributed computation
- Fault tolerance
- Efficient large-scale analytics

Innovation

An important innovation is the use of multiple intelligent agents instead of a single centralized decision-maker.

For example:

Performance Agent

+

Behavior Agent

+

Recommendation Agent

+

Feedback Agent

↓


Personalized Learning Plan

Conclusion

The Personalized Learning Analytics Platform combines PySpark RDDs, AI analytics, and intelligent agents to provide scalable and adaptive education. It continuously analyzes student behavior and performance and adjusts learning content according to individual needs.

Overall Data Flow for the 5 Systems

The common architecture can be represented as:



DECISION MAKING



USER / SYSTEM
RESPONSE



FEEDBACK



MODEL UPDATE

Important PySpark RDD Operations to Mention

Operation	Purpose
map()	Transform each data element
filter()	Select required records
flatMap()	Generate multiple outputs from one record
groupByKey()	Group values by key
reduceByKey()	Aggregate values efficiently
join()	Combine two RDDs
sortByKey()	Sort key-value data
count()	Count records
collect()	Retrieve RDD results
reduce()	Aggregate complete RDD